

# Big Data



## Big Data Engineering with Hadoop & Spark

Assignment on Sqoop



## Session 6: Assignment 6.1

---

This assignment is aimed at consolidating the concepts that was learnt during the Data Indigestion Tool Sqoop session of the course.

# Problem Statement

## Task 1:

- Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table.

Solution:

- Login to the “**mysql**” shell using the command  
`$ mysql -u root -p`
- Check databases available using the command  
`mysql> show databases;`
- From the list of the databases select “**simplidb**” database using the command  
`mysql> use simplidb;`

```
[acadgild@localhost ~]$ mysql -u root -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 8.0.3-rc-log MySQL Community Server (GPL)

Copyright (c) 2000, 2017, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show databases;
+-----+
| Database                |
+-----+
| information_schema      |
| metastore               |
| mysql                   |
| oozie                   |
| performance_schema      |
| simplidb                |
| sys                     |
+-----+
7 rows in set (0.04 sec)

mysql> use simplidb
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
```

- Check the tables available under “*simplidb*” database

```
mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| employee             |
+-----+
2 rows in set (0.00 sec)
```

- Create a table “*Person1*” to which data present in output folder needs to be exported to, use the command below

```
mysql> create table Person1(person_id INT NOT NULL, lname varchar(20),
fname varchar(20), area varchar(20), city varchar(20), PRIMARY KEY
(person_id));
```

```
mysql> Create table Person1(
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY (person_id)
-> );
Query OK, 0 rows affected (0.04 sec)

mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| Person1             |
| employee             |
+-----+
3 rows in set (0.00 sec)

mysql> desc Person1;
+-----+-----+-----+-----+-----+-----+
| Field      | Type          | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| person_id  | int(11)       | NO   | PRI | NULL    |       |
| lname      | varchar(20)   | YES  |     | NULL    |       |
| fname      | varchar(20)   | YES  |     | NULL    |       |
| area       | varchar(20)   | YES  |     | NULL    |       |
| city       | varchar(20)   | YES  |     | NULL    |       |
+-----+-----+-----+-----+-----+-----+
5 rows in set (0.01 sec)

mysql> █
```



- Execute the command to export data from output folder “*sqoopoutbyid*” to new table “*Person1*”

```
$ sqoop export --connect jdbc:mysql://localhost/simplidb --table Person1
--username root --password Root@123 --export-dir /sqoopoutbyid/ -m 1
--input-fields-terminated-by ',' --columns person_id,lname,fname,area,city
```

```
[acadgild@localhost ~]$
[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/simplidb --table Person1 --username root --password Root@123
--export-dir /sqoopoutbyid/ -m 1 --input-fields-terminated-by ',' --columns person_id,lname,fname,area,city
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/07/16 11:43:17 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
18/07/16 11:43:17 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
18/07/16 11:43:19 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/16 11:43:19 INFO tool.CodeGenTool: Beginning code generation
Mon Jul 16 11:43:21 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According
to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For
compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to expl
icly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/16 11:43:28 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person1' AS t LIMIT 1
18/07/16 11:43:28 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person1' AS t LIMIT 1
18/07/16 11:43:28 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/509c7f4e96c1229c839da3481d4a9d7f/Person1.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/07/16 11:43:48 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/509c7f4e96c1229c839da3481d4a9d7f/Perso
n1.jar
18/07/16 11:43:49 INFO mapreduce.ExportJobBase: Beginning export of Person1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/s
lf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLogger
Binder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/07/16 11:43:51 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes
where applicable
18/07/16 11:43:51 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/07/16 11:43:59 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.r
educe.speculative
18/07/16 11:44:00 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.
speculative
18/07/16 11:44:00 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
```

```
18/07/16 11:44:00 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/07/16 11:44:00 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/07/16 11:44:12 INFO input.FileInputFormat: Total input paths to process : 2
18/07/16 11:44:12 INFO input.FileInputFormat: Total input paths to process : 2
18/07/16 11:44:13 INFO mapreduce.JobSubmitter: number of splits:1
18/07/16 11:44:13 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.
speculative
18/07/16 11:44:15 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1531718966036_0001
18/07/16 11:44:20 INFO impl.YarnClientImpl: Submitted application application_1531718966036_0001
18/07/16 11:44:22 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1531718966036_0001/
18/07/16 11:44:22 INFO mapreduce.Job: Running job: job_1531718966036_0001
18/07/16 11:45:46 INFO mapreduce.Job: Job job_1531718966036_0001 running in uber mode : false
18/07/16 11:45:46 INFO mapreduce.Job: map 0% reduce 0%
18/07/16 11:46:43 INFO mapreduce.Job: map 100% reduce 0%
18/07/16 11:46:46 INFO mapreduce.Job: Job job_1531718966036_0001 completed successfully
18/07/16 11:46:47 INFO mapreduce.Job: Counters: 30
File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=127548
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=609
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=7
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=49626
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=49626
  Total vcore-milliseconds taken by all map tasks=49626
  Total megabyte-milliseconds taken by all map tasks=50817024
Map-Reduce Framework
  Map input records=12
  Map output records=12
```

```

FILE: Number of bytes read=0
FILE: Number of bytes written=127548
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=609
HDFS: Number of bytes written=0
HDFS: Number of read operations=7
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=49626
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=49626
  Total vcore-milliseconds taken by all map tasks=49626
  Total megabyte-milliseconds taken by all map tasks=50817024
Map-Reduce Framework
  Map input records=12
  Map output records=12
  Input split bytes=192
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=603
  CPU time spent (ms)=8830
  Physical memory (bytes) snapshot=119689216
  Virtual memory (bytes) snapshot=2061332480
  Total committed heap usage (bytes)=62980096
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/07/16 11:46:47 INFO mapreduce.ExportJobBase: Transferred 609 bytes in 167.1183 seconds (3.6441 bytes/sec)
18/07/16 11:46:47 INFO mapreduce.ExportJobBase: Exported 12 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

- Finally, to check if the data is exported in the **Person1** table, use the command below

```
mysql> select * from Person1;
```

```

mysql>
mysql> select * from Person1;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area       | city       |
+-----+-----+-----+-----+-----+
| 1         | Rizvi | Syed   | Whitefield | Bangalore  |
| 2         | Rizvi | Syed   | Whitefield | Bangalore  |
| 3         | Rizvi | Syed   | Whitefield | Bangalore  |
| 4         | Rizvi | Syed   | Whitefield | Bangalore  |
| 5         | Rizvi | Syed   | Whitefield | Bangalore  |
| 6         | Rizvi | Syed   | Whitefield | Bangalore  |
| 7         | Rizvi | Syed   | Whitefield | Bangalore  |
| 8         | Rizvi | Syed   | Whitefield | Bangalore  |
| 9         | Rizvi | Syed   | Whitefield | Bangalore  |
| 10        | Rizvi | Syed   | Whitefield | Bangalore  |
| 11        | Rizvi | Syed   | Whitefield | Bangalore  |
| 12        | Rizvi | Syed   | Whitefield | Bangalore  |
+-----+-----+-----+-----+-----+
12 rows in set (0.00 sec)

mysql> █

```



## Task 2:

- Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person\_id=3.

Solution:

- Import data with parameter “**person\_id=3**”, use the command below  
`$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root --password Root@123 --target-dir /sqoopoutbyPersonID -m 1 --where person_id=3`

```
[acadgild@localhost ~]$
[acadgild@localhost ~]$ sqoop import --connect jdbc:mysql://localhost/simplidb --table Person --username root --password Root@123 --target-dir /sqoopoutbyPersonID -m 1 --where person_id=3
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6/bin/hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6/bin/hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/07/16 13:49:15 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
18/07/16 13:49:15 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
18/07/16 13:49:16 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/07/16 13:49:16 INFO tool.CodeGenTool: Beginning code generation
Mon Jul 16 13:49:17 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/16 13:49:21 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/16 13:49:21 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person' AS t LIMIT 1
18/07/16 13:49:21 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acadgild/compile/7e2d491fdcec7f8fd9bea0lead407de8/Person.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/07/16 13:49:32 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/7e2d491fdcec7f8fd9bea0lead407de8/Person.jar
18/07/16 13:49:32 WARN manager.MySQLManager: It looks like you are importing from mysql.
18/07/16 13:49:32 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
18/07/16 13:49:32 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
18/07/16 13:49:32 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
18/07/16 13:49:32 INFO mapreduce.ImportJobBase: Beginning import of Person
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/07/16 13:49:33 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/07/16 13:49:33 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/07/16 13:49:36 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
```

```
18/07/16 13:49:37 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
Mon Jul 16 13:49:44 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/07/16 13:49:45 INFO db.DBInputFormat: Using read committed transaction isolation
18/07/16 13:49:45 INFO mapreduce.JobSubmitter: number of splits:1
18/07/16 13:49:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1531718966036_0004
18/07/16 13:49:48 INFO impl.YarnClientImpl: Submitted application application_1531718966036_0004
18/07/16 13:49:48 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1531718966036_0004/
18/07/16 13:49:48 INFO mapreduce.Job: Running job: job_1531718966036_0004
18/07/16 13:50:19 INFO mapreduce.Job: Job job_1531718966036_0004 running in uber mode : false
18/07/16 13:50:19 INFO mapreduce.Job: map 0% reduce 0%
18/07/16 13:50:46 INFO mapreduce.Job: map 100% reduce 0%
18/07/16 13:50:47 INFO mapreduce.Job: Job job_1531718966036_0004 completed successfully
18/07/16 13:50:48 INFO mapreduce.Job: Counters: 30
File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=128025
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=87
  HDFS: Number of bytes written=34
  HDFS: Number of read operations=4
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=22579
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=22579
  Total vcore-milliseconds taken by all map tasks=22579
  Total megabyte-milliseconds taken by all map tasks=23120896
Map-Reduce Framework
  Map input records=1
  Map output records=1
```

```

FILE: Number of bytes read=0
FILE: Number of bytes written=128025
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=34
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=22579
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=22579
  Total vcore-milliseconds taken by all map tasks=22579
  Total megabyte-milliseconds taken by all map tasks=23120896
Map-Reduce Framework
  Map input records=1
  Map output records=1
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=276
  CPU time spent (ms)=6190
  Physical memory (bytes) snapshot=121901056
  Virtual memory (bytes) snapshot=2063437824
  Total committed heap usage (bytes)=62980096
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=34
18/07/16 13:50:48 INFO mapreduce.ImportJobBase: Transferred 34 bytes in 71.7976 seconds (0.4736 bytes/sec)
18/07/16 13:50:48 INFO mapreduce.ImportJobBase: Retrieved 1 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

- Create a new table ***PersonExportByID*** to which data present in output folder ***sqoopoutbyPersonID*** needs to be exported to, use the command below

```
mysql> create table PersonExportByID(person_id INT NOT NULL, lname
varchar(20), fname varchar(20), area varchar(20), city varchar(20),
PRIMARY KEY (person_id));
```

```

mysql> Create table PersonExportByID(
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY (person_id)
-> );
Query OK, 0 rows affected (0.03 sec)

mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| Person1             |
| PersonExportByID    |
| employee            |
+-----+
4 rows in set (0.00 sec)

```



- Execute the command to export data from output folder “*sqoopoutbyPersonId*” to new table “*PersonExportByID*”, use the command below

```
$ sqoop export --connect jdbc:mysql://localhost/simplidb --table  
PersonExportByID --username root --password Root@123 --export-dir  
/sqoopoutbyPersonID/ -m 1
```

```
[acadgild@localhost ~]$  
[acadgild@localhost ~]$ sqoop export --connect jdbc:mysql://localhost/simplidb --table PersonExportByID --username root --password  
Root@123 --export-dir /sqoopoutbyPersonID/ -m 1  
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./hcatalog does not exist! HCatalog jobs will fail.  
Please set $HCAT_HOME to the root of your HCatalog installation.  
Warning: /home/acadgild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
18/07/16 14:00:01 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6  
18/07/16 14:00:01 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
18/07/16 14:00:02 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
18/07/16 14:00:02 INFO tool.CodeGenTool: Beginning code generation  
Mon Jul 16 14:00:03 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According  
to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For  
compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to expl  
icly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.  
18/07/16 14:00:08 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'PersonExportByID' AS t LIMIT 1  
18/07/16 14:00:08 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'PersonExportByID' AS t LIMIT 1  
18/07/16 14:00:08 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acadgild/install/hadoop/hadoop-2.6.5  
Note: /tmp/sqoop-acadgild/compile/622cf19a1e3cd9af37ed59125b36c27d/PersonExportByID.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
18/07/16 14:00:18 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acadgild/compile/622cf19a1e3cd9af37ed59125b36c27d/Per  
sonExportByID.jar  
18/07/16 14:00:18 INFO mapreduce.ExportJobBase: Beginning export of PersonExportByID  
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/s  
lf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLogger  
Binder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]  
18/07/16 14:00:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes  
where applicable  
18/07/16 14:00:19 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
18/07/16 14:00:24 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.r  
educe.speculative  
18/07/16 14:00:24 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.  
speculative  
18/07/16 14:00:24 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
```

```
18/07/16 14:00:24 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
18/07/16 14:00:25 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032  
18/07/16 14:00:31 INFO input.FileInputFormat: Total input paths to process : 1  
18/07/16 14:00:31 INFO input.FileInputFormat: Total input paths to process : 1  
18/07/16 14:00:32 INFO mapreduce.JobSubmitter: number of splits:1  
18/07/16 14:00:32 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.  
speculative  
18/07/16 14:00:33 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1531718966036_0005  
18/07/16 14:00:35 INFO impl.YarnClientImpl: Submitted application application_1531718966036_0005  
18/07/16 14:00:35 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1531718966036_0005/  
18/07/16 14:00:35 INFO mapreduce.Job: Running job: job_1531718966036_0005  
18/07/16 14:01:06 INFO mapreduce.Job: Job job_1531718966036_0005 running in uber mode : false  
18/07/16 14:01:06 INFO mapreduce.Job: map 0% reduce 0%  
18/07/16 14:01:37 INFO mapreduce.Job: map 100% reduce 0%  
18/07/16 14:01:38 INFO mapreduce.Job: Job job_1531718966036_0005 completed successfully  
18/07/16 14:01:39 INFO mapreduce.Job: Counters: 30  
File System Counters  
FILE: Number of bytes read=0  
FILE: Number of bytes written=127581  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=171  
HDFS: Number of bytes written=0  
HDFS: Number of read operations=4  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=0  
Job Counters  
Launched map tasks=1  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=25294  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=25294  
Total vcore-milliseconds taken by all map tasks=25294  
Total megabyte-milliseconds taken by all map tasks=25901056  
Map-Reduce Framework  
Map input records=1  
Map output records=1
```

```

FILE: Number of bytes read=0
FILE: Number of bytes written=127581
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=171
HDFS: Number of bytes written=0
HDFS: Number of read operations=4
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=25294
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=25294
  Total vcore-milliseconds taken by all map tasks=25294
  Total megabyte-milliseconds taken by all map tasks=25901056
Map-Reduce Framework
  Map input records=1
  Map output records=1
  Input split bytes=134
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=287
  CPU time spent (ms)=5540
  Physical memory (bytes) snapshot=118603776
  Virtual memory (bytes) snapshot=2061332480
  Total committed heap usage (bytes)=62980096
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/07/16 14:01:39 INFO mapreduce.ExportJobBase: Transferred 171 bytes in 74.3469 seconds (2.3 bytes/sec)
18/07/16 14:01:39 INFO mapreduce.ExportJobBase: Exported 1 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

- Finally, to check if the data is exported in the “*PersonExportByID*” table, use the command below  
*mysql> select \* from PersonExportByID;*

```

mysql>
mysql> select * from PersonExportByID;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area       | city       |
+-----+-----+-----+-----+-----+
|          3 | Rizvi | Syed   | Whitefield | Bangalore  |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql>

```